

Kalman filter tutorial for beam current tracking

Victor Watson¹

¹Nuclear Science Division
Lawrence Berkeley National Laboratory

June 2025

Introduction

Kalman Filter

- Widely used for **tracking** purposes
- Can be seen as an online optimal **low pass filter**
- Has an **adaptive gain** based on **Bayesian statistics**

Characteristics

- As a filter it is a **linear** and **invariant system**
- Still behaves **very well with non-linear** signals
- Gives a **prediction** of the current state based on a **linearisation** of past data
- **Updates** the prediction with the **measure** and gives the **estimate** thanks to the **optimal gain**

Theory

- The prediction of the state is given by:

$$X_p(k) = \hat{X}(k-1) + \frac{dX}{dt} * \delta t + \mathcal{N}(0, \mathbf{P}_p(k)) \quad (1)$$

with $\hat{X}(k-1)$ the previous estimated state and \mathbf{P}_p the prediction covariance matrix

- The measure follows:

$$Z_m(k) = Z(k) + \mathcal{N}(0, \mathbf{R}) \quad (2)$$

with $Z(k)$ the real value and \mathbf{R} the covariance of the noise

- The estimate is:

$$\hat{X}(k) = X_p(k) + (\mathbf{K} * Z_m(k))^T \quad (3)$$

Theory

- With the optimal Kalman gain:

$$\mathbf{K} = (\mathbf{P}(k-1)\mathbf{H}^T)^T \mathbf{S}(k)^{-1} \quad (4)$$

\mathbf{H} is the transition matrix from state space to measure space

- The prediction covariance matrix is calculated:

$$\mathbf{P}_p(k) = \mathbf{F} * \mathbf{P}(k-1)\mathbf{F}^T + \mathbf{Q} \quad (5)$$

where F is the state transition matrix \mathbf{P} is the estimate covariance matrix and \mathbf{Q} is a matrix set by the user inversely proportional to the confidence they have that the system is linearisable.

Theory

- **S** is the covariance of the innovation:

$$\mathbf{S}(k) = \mathbf{H} * \mathbf{P}_p(k) + \mathbf{R} \quad (6)$$

where **R** is set by the user and is proportional to the standard deviation of the measure noise

Algorithm

- Initialization
 - Set **F**, **H**, **R**, **Q**, $\hat{X}(0)$ (we use the same 4 matrices for $\hat{\sigma}_{beam}$)
 - **F** is the transition matrix from a state to the next if $X = [position, \frac{dx}{dt}]^T$ then $F = [[1, \delta t], [1, 0]]$ that way $X(k) = [X(k-1)[0] + \frac{dx}{dt} \delta t, \frac{dx}{dt}]^T$ we suppose the system locally linear so **F** propagate dx .
 - **H** usually the first part of X (in our case $X(k)[0]$) is homogeneous to the measure so $\mathbf{H} = [1, 0]^T$ (state is 2d measure is 1d)

Algorithm

- Initialization
 - $\hat{X}(0)$ **best guess** about the state of the space in the beginning usually the first measure for $X(0)[0]$ and start with a null dynamic (unless better idea) $X(0)[1] = 0$
 - **R** Covariance of the measure dimension of the measure (in our case a scalar) pick a **big number** because we assume the measure is imprecise
 - **Q** Covariance of the prediction dimension of the state (in our case 2x2) pick a **small number** because we assume the system is mostly linearisable (eye(2)*something small) the **ratio R/Q** will give the reactivity or the cutting frequency of the filter if this ratio is **too small**, the filter will follow the noise and be completely **useless**, if it is **too big** the filter will take a **longer time** to catch-up when the **trajectory** of the system changes

Prediction

- Prediction

- Predicting X :

$$X_p(k) = \mathbf{F}\hat{X}(k) \quad (7)$$

- Covariance of predicted X :

$$\mathbf{P}_p(k) = \mathbf{F}\mathbf{P}(k-1)\mathbf{F}^T + \mathbf{Q} \quad (8)$$

- Predicting σ :

$$\sigma_p(k) = \mathbf{F}\hat{\sigma}(k) \quad (9)$$

- Covariance of predicted X :

$$\mathbf{P}_{s_p}(k) = \mathbf{F}\mathbf{P}_s(k-1)\mathbf{F}^T + \mathbf{Q} \quad (10)$$

Theory and Algorithm

Estimation

- Update

- Innovation of X (z is the measure):

$$y = z(k) - \mathbf{H}X_p(k) \quad (11)$$

- Covariance of innovation:

$$\mathbf{S}(k) = \mathbf{H}\mathbf{P}_p(k)\mathbf{H}^T + \mathbf{R} \quad (12)$$

- Optimal Kalman Gain:

$$\mathbf{K}(k) = (\mathbf{P}(k-1)\mathbf{H}^T)^T \mathbf{S}(k)^{-1} \quad (13)$$

- Estimating X

$$\hat{X}(k) = X_p(k) + (\mathbf{K}(k)y)^T \quad (14)$$

- updating X covariance:

$$\mathbf{P}(k) = (\mathbb{1}_{2 \times 2} - \mathbf{K}\mathbf{H})\mathbf{P}_p(k) \quad (15)$$

Estimation

- Update

- Innovation of σ using \hat{X} :

$$y = \sqrt{(z(k) - \hat{X}(k)[0])^2 - \mathbf{H}\sigma_p(k)} \quad (16)$$

- Covariance of innovation:

$$\mathbf{Ss}(k) = \mathbf{H}\mathbf{Ps}_p(k)\mathbf{H}^T + \mathbf{R} \quad (17)$$

- Optimal Kalman Gain:

$$\mathbf{Ks}(k) = (\mathbf{Ps}(k-1)\mathbf{H}^T)^T \mathbf{Ss}(k)^{-1} \quad (18)$$

- Estimating σ

$$\hat{\sigma}(k) = \sigma_p(k) + (\mathbf{Ks}(k)y)^T \quad (19)$$

- updating σ covariance:

$$\mathbf{Ps}(k) = (\mathbb{1}_{2 \times 2} - \mathbf{Ks}(k)\mathbf{H})\mathbf{Ps}_p(k) \quad (20)$$

object KFpar

- Contains all the **variables** necessary to characterize the state and **parameters** of the filter
- Is initialized with **cLin** for **Q** and **cMeas** for **R**^a
- Receive the first measure in **KFpar.X[0]** and the elapsed time before every update in **KFpar.F[0,1]**
- Store Current KFpar.X[0], slope KFpar.X[1] and stdev KFpar.Sig[0]

^acLin and cMeas are scalar used to set **Q** and **R** considering $\mathbf{Q} = \mathbb{1}_{StateDim} * cLin$ and $\mathbf{R} = \mathbb{1}_{MeasureDim} * cMeas$

function EstimateState

- Updates the estimate of the Kalman Filter and its parameters with the new measure (KFpar.F needs to be updated before)
- **KFpar = EstimateState(KFpar, measure)**

Examples

Data

- Series of examples generated with the **KFtutorial.py** and **KFtutorialAtan.py**
- These are **illustrating** the impact of the **filter settings** in several cases
- These codes are adapted to track a 1D signals and extract from it estimates of the mean the slope and the standard deviation

Cases studies

- KFtutorial.py applies the Kalman filter to one of the times series in the files 'd1', 'd2', or 'd3'
- KFtutorialAtan.py allows the user to change the settings of an arctangent and add perturbation to it before applying the Kalman Filter

Example 'd1'

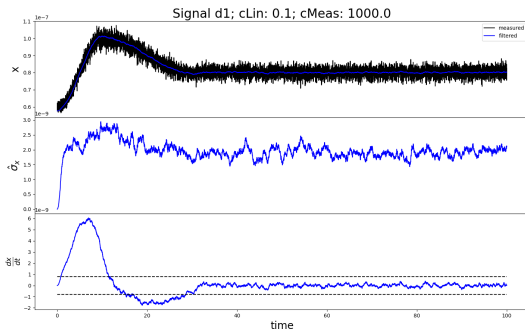


Figure: $R = 1000$; $Q = 0.1$

Q and R are adapted to most of the noise is filtered but x is tracked successfully (Note: the estimate of σ is impacted by the slope in the beginning)

Example 'd1'




Figure: $R = 1$; $Q = 0.1$

Q and R are not adapted and **too much noise** is still impacting the estimate

Example 'd1'




Figure: $R = 100$; $Q = 1000$

Q and R are not adapted and **too much noise** is still impacting the estimate note result is similar to the previous figure because ratio Q/R is the same

Example 'd1'

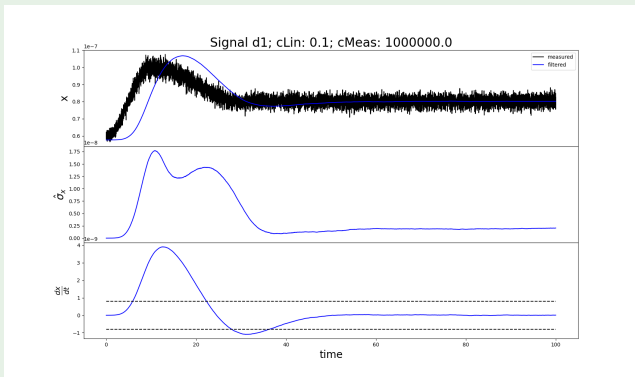


Figure: $R = 1000000$; $Q = .1$

Q and R are not adapted and the filter **cannot track** correctly the changes in x the delay impacts the slope estimate even more

Example 'd2'

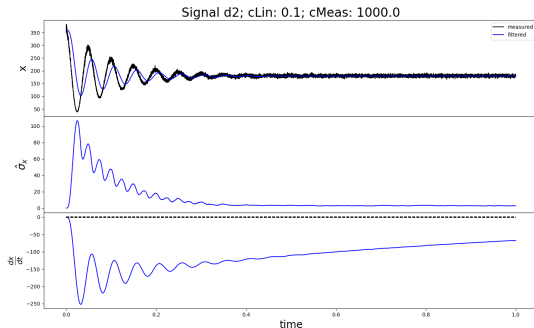


Figure: $R = 1000$; $Q = .1$

Q and R are not adapted and the filter **cannot track** correctly the changes in x

Example 'd2'

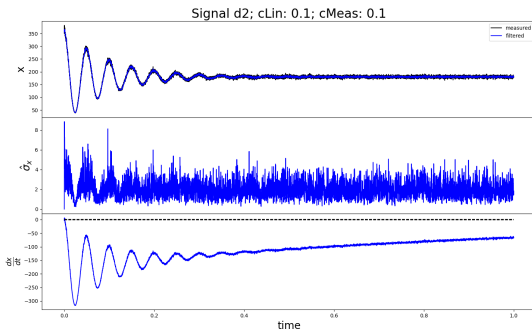


Figure: $R = .1$; $Q = .1$

Q and R are adapted to track x but the filter is **not very robust to noise** and the **slope convergence** is way **too slow**. In this case we see the limits of the assumption of the linearisable aspect of the signal.

Example 'd3'




Figure: $R = 1000$; $Q = .1$

Q and R are adapted to track x and filter most of the noise but takes longer to estimate the slope due to the change in direction.

Example Atan




Figure: $R = 1000$; $Q = .1$

Q and R are adapted to track x and filter most of the noise.

Example Atan

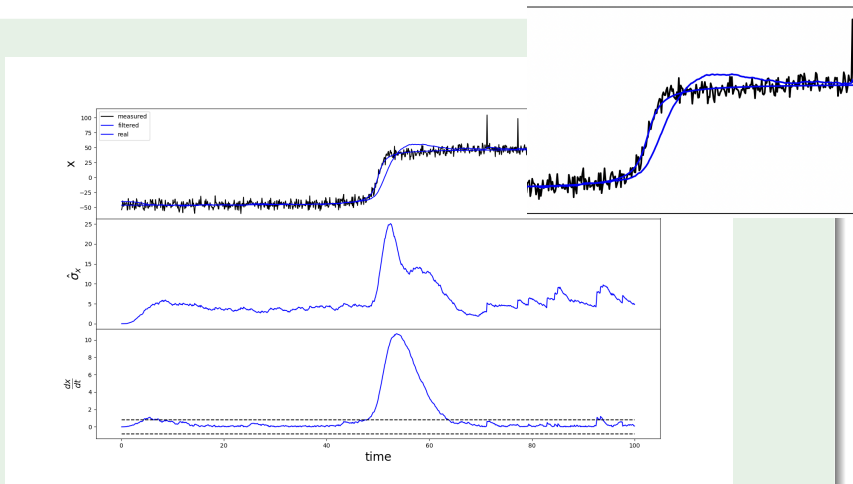


Figure: $R = 1000$; $Q = .1$

Q and R are not adapted to track x . Here the settings did not change but the change of **trajectory is sharper** and the settings are no longer adapted

Conclusion

Essential

- This Kalman filter is **adapted to a 1D signal** and estimate the measured value as well as the slope and standard deviation
- The use of this tool is valid as long as the system dynamic is **sufficiently linearisable vis a vis noise**
- If **brutal changes** in direction add detection and **reset filter**

How to

Set **cLin** (small) and **cMeas** (big) to determine the reactivity vs cutting frequency adapted to the dynamic and the noise

initialize KFpar.X[0] with the first measure

At every measure

- Enter time to since the last measure **$F[0,1] = \delta t$**
- **Update the Object:** **KFpar = EstimateState(KFpar, measure)**

Current: **KFpar.X[0]** ; Slope: **KFpar.X[1]** ; Stdev: **KFpar.Sig[0]**

Compact alternative

- AlternateKFtutorial.py contains a **compact alternative** that does the same thing
- It encapsulate the function EstimateState as a method of the object **KFparAlt**

How to

Set **cLin** (small) and **cMeas** (big) to determine the reactivity vs cutting frequency adapted to the dynamic and the noise (ex: `myKFparAlt = KFparAlt(.1,1000)`)

initialize `myKFparAlt.X[0]` with the first measure

At every measure

- **Update the Object:** `myKFparAlt.EstimateState(measure, deltaT)`

Current: `myKFparAlt.X[0]` ; Slope: `myKFparAlt.X[1]` ; Stdev: `myKFparAlt.Sig[0]`